

클라우드 기반 빅데이터 환경구축 (Part 2)

실습조교: 이계원, 엄태건
서울대학교 컴퓨터공학부

Overview

0. 실습에서 사용할 소프트웨어
1. 빅데이터 분석을 위한 클라우드 환경 셋업
2. 클라우드상 배치 처리 분석
3. 클라우드상 대화형 질의 분석
- 4. 클라우드상 스트림 처리 분석**
- 5. 클라우드상 기계학습 분석**
6. 클라우드상 딥러닝 분석

4. 클라우드상 스트림 처리 분석

Spark Streaming Basics (Review)

Stream Processing

- Processes unbounded incoming stream data in real-time
- Example: Real-time network log aggregation, Twitter trend analysis, Anomaly detection, ...
- vs Batch processing
 - Batch processing processes **bounded** set of data and produces results
 - Stream processing processes **unbounded** incoming stream data events **continuously**
- Two types of execution models
 - Continuous operator model
 - Micro-batch model


Spark Streaming Basic

- Spark Streaming Implement micro-batch model stream processing on top of Apache Spark
- Run a streaming computation as a series of small, deterministic batch jobs (**micro-batches**)




Spark Streaming Query Processing Example

Stream Events




**Micro-batches
(batch size = 5)**



`.map(lambda x: x * 2)`

Processed data



Spark Streaming Basic (Contd.)

- vs Continuous operator model (Storm, Heron)
 - Advantage
 - Fast fault recovery via parallel recovery
 - Straggler handling
 - Disadvantage
 - High latency (second-scale) caused by micro-batch processing & scheduling overhead

Stream Processing via Spark Streaming



Spark Streaming Practice

Prerequisites

- Java 8 or higher version
- AWS YARN cluster with HDFS setup
- Spark cluster working on YARN


What we will do

- Part 1: Implementing & running simple wordcount example query
- Part 2: Setting up Kafka source for large-scale message transportation
- Part 3: Implementing & running windowed movie aggregation

Detailed information can be found on Jupyter notebook

Part 1: Running a simple wordcount query

- Continuously get the text stream data via TCP socket
- Split the sentence into multiple words
- Update the word count data for each word
- Emit the output to a text file
- Test the running query with netcat



Part 1 Detail (1) - Set up streaming context

```
spark_context = SparkContext(...)
```

```
streaming_context = StreamingContext(sc, 1)
```



**Batch interval
(1 second)**

Part 1 Detail (2) - Write stream query

Fetch the line from a tcp socket

```
lines = streaming_context.socketTextStream(...)
```

Processing something

...

Output the result to the save

```
results.saveAsTextFiles(...)
```

Part 2: Setting Up Kafka source

- Download the movie data json set (<https://raw.githubusercontent.com/prust/wikipedia-movie-data/master/movies.json>)
- Install & configure Kafka messaging server
- Make a script which publishes real-time movie watching data stream to the Kafka messaging server

Apache Kafka

- Kafka is a distributed messaging & streaming platform
 - We will use messaging features of Kafka in this tutorial
- Kafka supports large-scale & reliable publishing / subscribing data stream events per topic
 - Similar to message brokers (RabbitMQ, ...)
- Input data is partitioned, distributed, and replicated to multiple machines to guarantee fault tolerance
- For more information, consult to
 - <https://kafka.apache.org/documentation/#introduction>

Part 3: Windowed Movie Popularity Aggregation

- Get the json-typed movie data stream from Kafka
- Parse the json data into key-value data stream
- Implement windowed aggregation
- Write the result to HDFS file

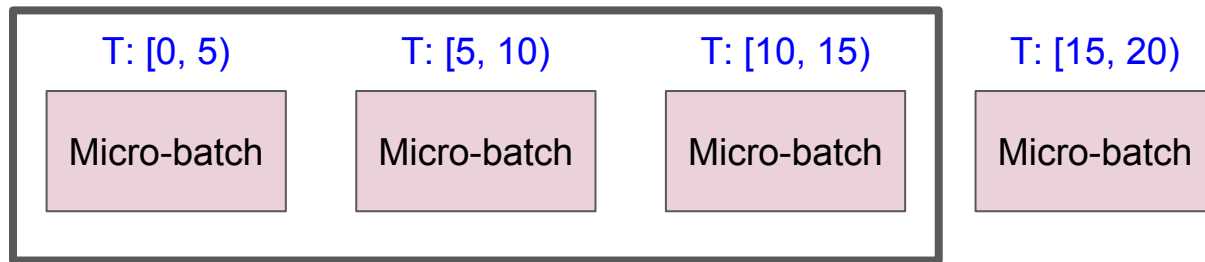
Recap: Window Operation

- Spark Streaming allows processing data over a recent set of data, called **Window**
- The types of windows are
 - Time-based window
 - Count-based window
 - Session window
- We will implement time-based window
 - Window Length & Slide Interval

Sliding Window Operation in Spark Streaming

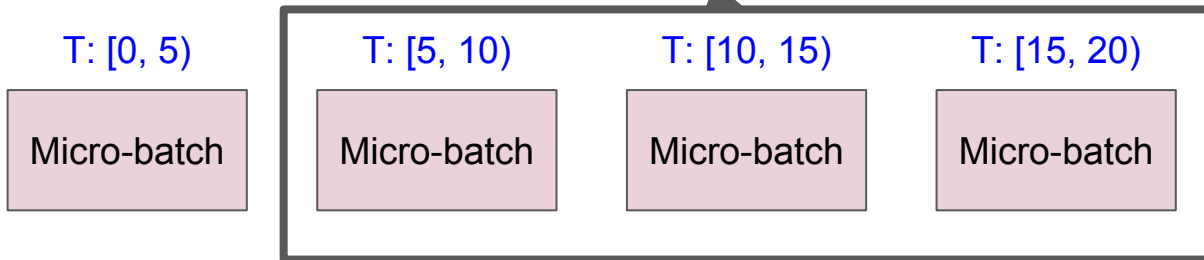
Window Size = 20, Slide Interval = 5

1st Window
[0, 15)



Sliding!

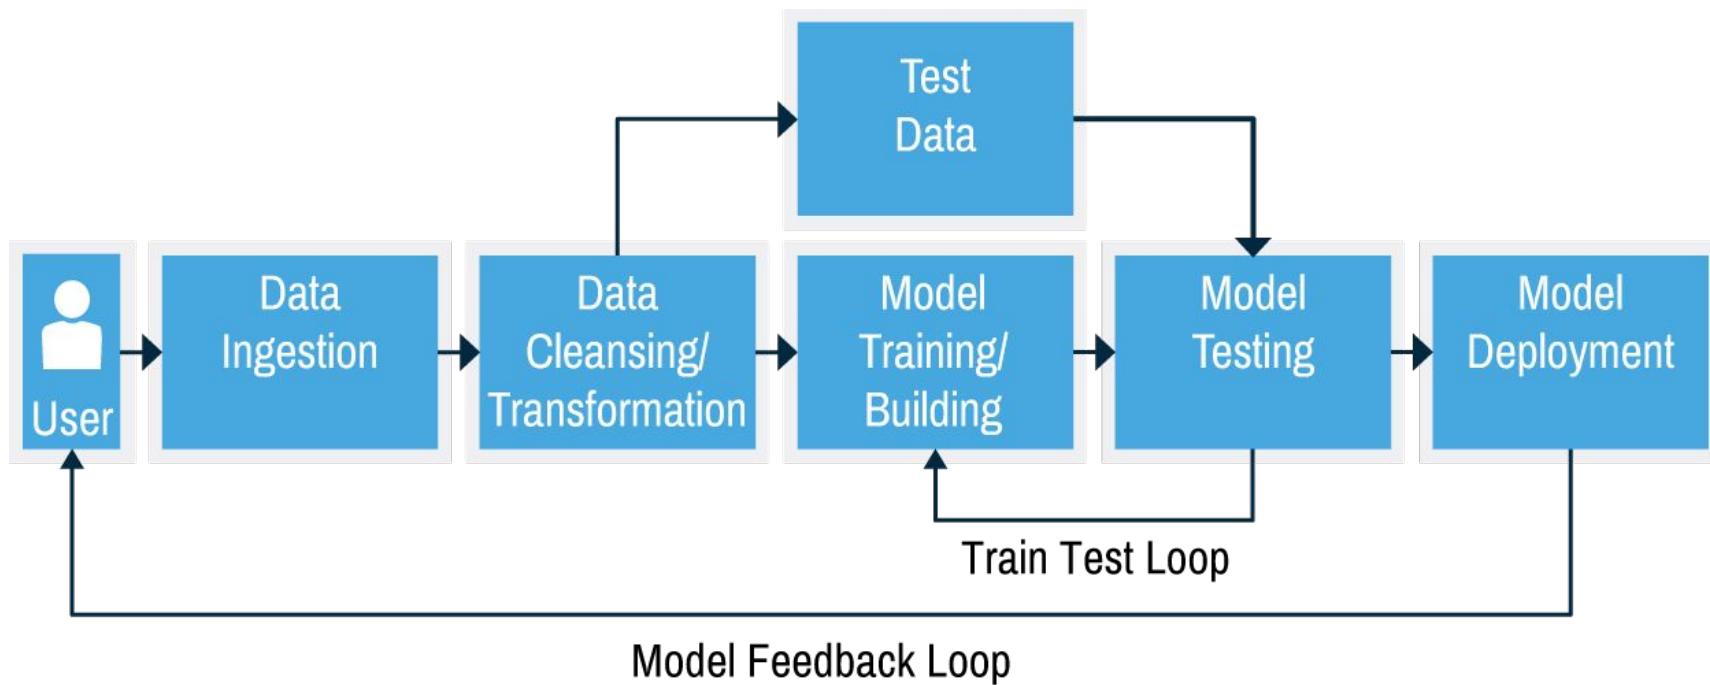
2nd Window [5, 20)



5. 클라우드상 기계학습 분석

Spark MLlib을 활용한 영화 추천 모델 만들기

	Movie 1	Movie 2	Movie 3
Ted	4	5	5
Carol		5	5
Bob		5	?

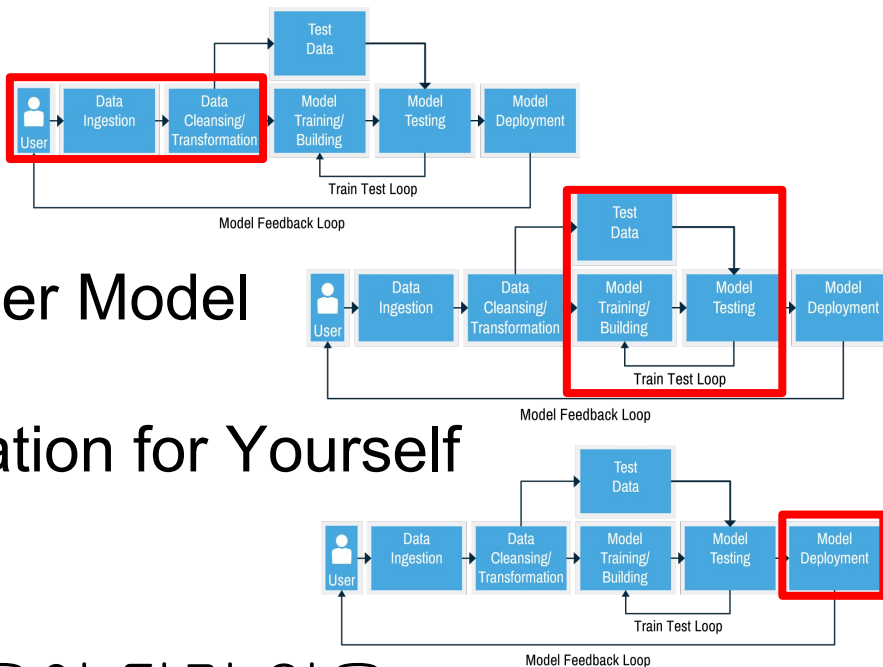


Typical Machine Learning Workflow

Movie Recommendation with Spark MLlib (ALS)

Contents Overview

- Part 1: Preparing your data
- Part 2: Build a Recommender Model
- Part 3: Movie Recommendation for Yourself



Jupyter notebook에 자세한 내용이 담겨 있음

Part 1: Preparing your data

We will use the following dataset:

1. ../data/ratings.data.gz

a. <MovieName>::<Rating>::UserId

```
One Flew Over the Cuckoo's Nest (1975)::5::1  
James and the Giant Peach (1996)::3::1  
My Fair Lady (1964)::3::1
```

2. ../data/movies.dat


a. <MovieId>::<MovieName>::Genre

```
1::Toy Story (1995)::Animation|Children's|Comedy  
2::Jumanji (1995)::Adventure|Children's|Fantasy  
3::Grumpier Old Men (1995)::Comedy|Romanc
```

Part 1: Preparing your data

One Flew Over the Cuckoo's Nest
(1975)::5::1
James and the Giant Peach (1996)::3::1
My Fair Lady (1964)::3::1

1::Toy Story
(1995)::Animation|Children's|Comedy
2::Jumanji
(1995)::Adventure|Children's|Fantasy
3::Grumpier Old Men
(1995)::Comedy|Romanc



DataFrame

UserId	Movielid	Rating
1	1	5
2	5	3
1	3	3

Part 2: Build a Movie Recommender Model

1. Create an ALS object and build a recommender model

```
from pyspark.ml.recommendation import ALS
```

```
# Create an ALS object
```

```
als = ALS(rank=8, maxIter=5, numUserBlocks=10, numItemBlocks=10,  
          implicitPrefs=False, userCol='userId',  
          itemCol='movieId', ratingCol='rating',  
          coldStartStrategy="drop")
```

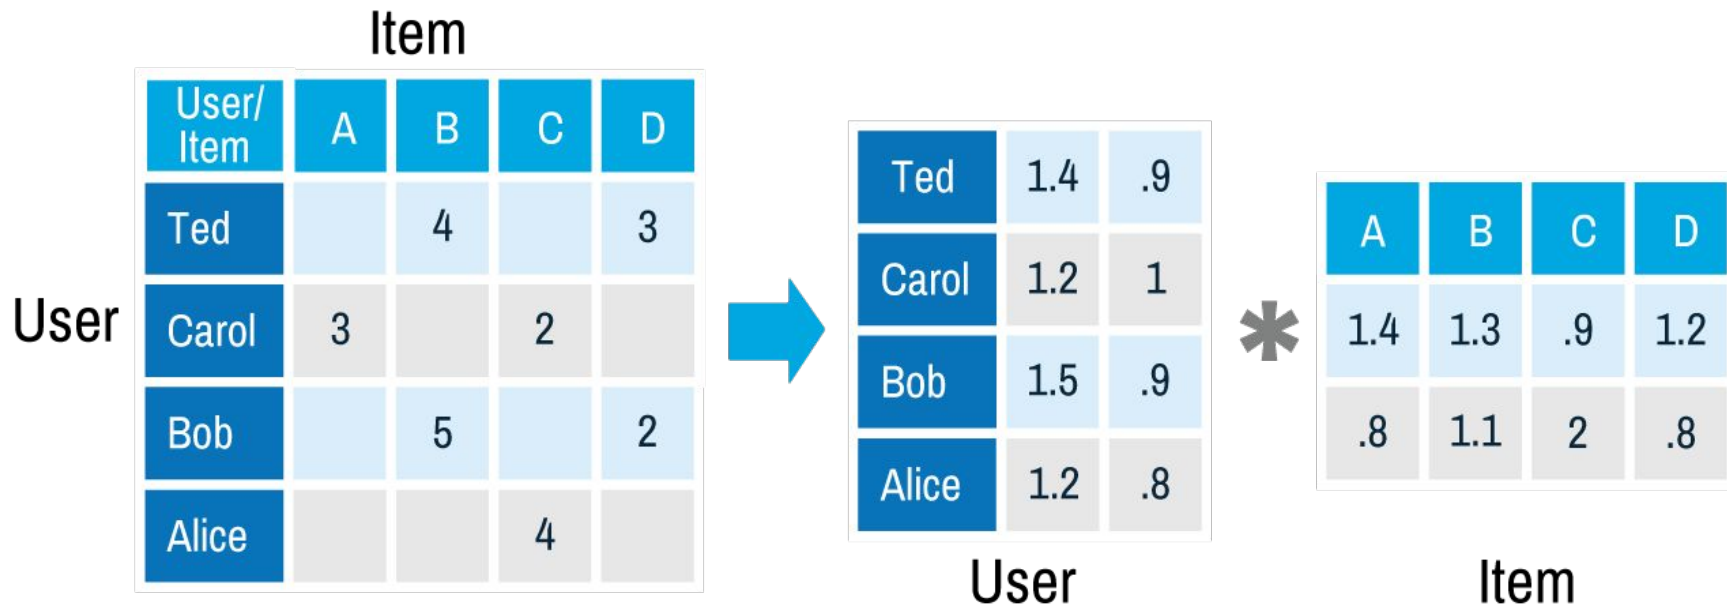
userId	movieId	rating
1	1	5
2	5	3
1	3	3

```
# Build our first recommender model!
```

```
# Spark automatically trains the model with your dataset
```

```
firstModel = als.fit(ratings)
```

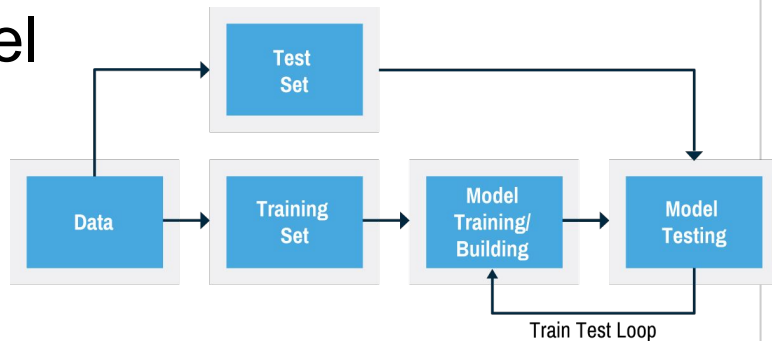
Part 2: Build a Movie Recommender Model



Part 2: Build a Movie Recommender Model

2. Improve your recommender model

- a. Split your dataset!
 - i. Training set
 - ii. Test set



```
# Split the dataset into 80% and 20% testSet
trainingSet, testSet = (ratings.randomSplit([0.8, 0.2]))
```

Part 2: Build a Movie Recommender Model

2b. Estimating error using Root Mean Square Error (RMSE)

```
from pyspark.ml.evaluation import RegressionEvaluator
```

```
evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",  
predictionCol="prediction")
```

```
rmse = evaluator.evaluate(testPrediction)
```

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (\text{Predicted}_i - \text{Actual}_i)^2}{N}}$$

2c. Let's try different parameters and choose the set of parameters which results in lowest RMSE value.

Part 3: Movie Recommendation for Yourself

- 3a. Make your dataset of movie ratings
- 3b. Train the recommender model with your dataset
- 3c. Recommend movies for you by using
`recommendForUserSubset` method

```
# Pick 10 movie recommendations for you
myRatingPrediction = model.recommendForUserSubset(myRatingsSet, 10)

myRatingPrediction.show()
print(myRatingPrediction.select("recommendations").take(10))
```

Homework #1:

AWS, HDFS, Spark SQL, Spark Streaming

Homework #1-1: AWS, HDFS, Spark SQL

- You'll do by yourself...
 - Launch AWS instances
 - Load data into HDFS
 - Run a Spark SQL query, and visualize the results
- You'll submit...
 - Screenshots of Spark Web UI
 - Screenshot of visualized output dataset

Homework #1-2: Stream Join

- You'll do yourself...
 - Generate KAFKA datastream
 - Transform the input streams into keyed window
 - Perform join operation on data streams
 - Write the result to HDFS
- You'll submit...
 - The stream query source code
 - The 1-min running result file written to HDFS